30-day development plan for the Associate

Part D(i)(c). Goal: in 30 days the Associate can take a change from a ticket to a released, measured result on this system without rescue, and has the habit of reporting honest numbers. The review (D1) shows a capable builder with good structure and three gaps: **measure honestly, treat safety rules as code, and respect ownership boundaries** (a service, a key, a number). The plan is built around those, not around new tools.

How I will work with them. A 30-minute 1:1 twice a week (Tuesday: unblock and review; Friday: look at the week's numbers together). Pair for the first hour of each new skill. PRs reviewed within one working day, with comments that explain why. I will say what I am doing and why, so the Associate can do the same later. Success is theirs to define each week in writing, and my feedback measures against that.

Week	Focus	Concrete work	Evidence it is done
1: Fix and understand (days 1 to 7)	Own the review findings	Restore `stub_tools.py` and move the key (day 1). Fix the missing-input and hard-gate defects with tests that need no API (D1 B3, B4). Write a short note: how did the key, the stub edit and the 95% line happen, and what will you check next time?	The tests pass and fail in the right cases. The note exists and is honest. Walk me through the gate code and explain every branch.
2: Measure (days 8 to 14)	Evaluation habits	Build a 15 to 20 question evaluation set for their own version, written from the documents before running anything. Compute scores in code; report failures. Learn from my `eval/` folder: labelled data, deterministic scoring and a judge for free text, and why each. Read 5 failed answers by hand.	A report with every failure listed. A one-page "what I trust and do not trust about this score". No hard-coded numbers anywhere.
3: Retrieval quality (days 15 to 21)	Chunking, filters, ambiguity	Replace the fixed 500-character chunks with heading-based chunks; add a registry with several categories per document; make the category filter soft with a fallback (D1 H3 to H5). Measure recall@k before and after.	A before/after table on their own set. They can explain, without notes, why document 06 is handled as it is and one risk of their approach.
4: Ship a change (days 22 to 30)	Working in the team's process	Take a real ticket from the backlog (for example the conflict-detection item from the TDR, or the cost-per-trace report). Follow the SOP in C1 and the change control in C2 end to end: branch, evals, comparison table, approval, release, 7-day check. Present it at the team review.	The change is released with a comparison table and a rollback target. Their PR needs at most one round of comments.

What I will check at day 30

- **Safety by construction:** every rule that must always hold (the gate, no key in code) is a test, not a hope.
- **Honest numbers:** every figure in their work traces back to a report the code generated, with caveats.
- **Boundaries:** they ask the owner before editing something that belongs to someone else.
- **Communication:** a PR description saying what, why, how tested, and what could go wrong.
- **Independence:** they pick the next ticket's approach themselves and ask questions early.

If it is going slower

At day 10 and day 20 I check progress against the table. If the tests and measurement habits are not in place by day 14, week 3 is shortened and the retrieval work moves to the following month; the safety and measurement habits come first, since the review shows those are the gaps that matter. I will say that plainly to the Associate at the time, not at day 30.

What I own

Clear tickets with acceptance criteria. Review within a day. A weekly honest summary. Protecting 3 uninterrupted hours a day for this work. Making sure this plan does not collide with the delivery plan (D4): the tasks in weeks 3 and 4 are chosen to be the same tasks the delivery plan needs.